Six-member speaking guide

What to say, what to show, and professor questions

Aligned with the simplified 18-slide presentation. 20 minutes presenting, followed by 5 minutes of questions.

How to use this guide

Each member has three slides and approximately 3 minutes 20 seconds, including pointing to diagrams and allowing the audience to inspect examples. Rehearse at your natural pace. These are speaking assignments, not claims about who completed the underlying work.

Introduce everyone by name. Each member should add one short, truthful sentence about their actual contribution. Keep the main presentation in the slides; use browser evidence or backup slides during questions.

The story

Client need to comparison design to page changes to measured results to interpretation to handover.

Speaking order

Member 1: slides 1-3 | Problem and system

Member 2: slides 4-6 | Comparison and method selection

Member 3: slides 7-9 | Implementation and page changes

Member 4: slides 10-12 | Checks, demonstration and results

Member 5: slides 13-15 | Confidence and practical meaning

Member 6: slides 16-18 | Limitations and handover

Answering questions

Answer directly, show one supporting piece of evidence, then explain the boundary. If the experiment does not establish something, say: "We have not established that from this experiment. What we can show is..."

AI assistance

Acknowledge AI assistance accurately and retain authorship records. Team members must verify and understand the material they present. The handover agreement still requires confirmation of the recipient, date, channel and acceptance process.

Member 1 | Problem and system

Slides 1-3 | About 3:20 total, including visual pauses

Slide 1: Making Flight Centre information easier for AI to use

SHOW: Point to the central question. Briefly introduce the six members by name.

We are presenting our Semester Two work for Flight Centre Travel Group. Each member will briefly introduce themselves as their section begins. Flight Centre wanted to understand how its travel products could become more visible and relevant in AI-assisted shopping. We built and tested a page-improvement prototype. Our main question is simple: when an AI already has Flight Centre information alongside competitors, does a clearer page make it more likely to cite Flight Centre? We will show the system, the page changes, a saved demonstration and the results. The talk is twenty minutes, followed by five minutes for questions.

Slide 2: What we measured

SHOW: Trace Find, Cite, Recommend and Book. Pause on Cite: that is the main measurement.

Imagine an answer that gives travel information and places a small source number next to it. If that number points to a Flight Centre page, we count the answer as citing Flight Centre. We count the answer once, even if it includes several Flight Centre references. This is one part of the larger customer journey. Finding a page, citing it, recommending the business and making a booking are different outcomes. This semester's main result concerns the citation step. It gives evidence about being used as an information source, not proof of increased sales.

Member 1 | Continue and hand over

Slide 3: How our system works

SHOW: Trace Search, Save, Ask AI and Check. Point out who controls the reading pack and who chooses the references.

Think of our system as giving the AI a reading pack. First, our search code finds travel pages. We save copies so we know which information was used. Then the system prepares a question and numbered page text for the AI. The model receives that information through an API and writes an answer with source references. Our checker reads those references and identifies whether Flight Centre was cited. This is often called retrieval-augmented generation, or RAG: supplying information to help a model answer. We did not train a new model. During the main answer test, the AI did not browse the live web or choose new pages outside the supplied pack.

HANDOVER: The next section explains how we compared the original and improved pages.

Questions you can lead

Did you force the AI to choose Flight Centre?

No. We supplied Flight Centre and competitor sources, but the model chose its references. We studied cases where Flight Centre was already available, so this does not measure unrestricted discovery.

Does a citation mean a booking recommendation?

No. It means the answer references a Flight Centre source. Recommending Flight Centre as a booking provider is a separate outcome.

What is RAG here?

It means supplying saved source information to help the model answer. We did not train a new model or run live web browsing during the main generation test.

Member 2 | Comparison and method selection

Slides 4-6 | About 3:20 total, including visual pauses

Slide 4: A fair before-and-after comparison

SHOW: Compare the Before and After branches. Point to the same competitor pages on each side.

The diagram shows two reading packs. One contains the baseline Flight Centre page and the other its improved representation. We match the question, source identities, source order, competitor text and instruction. We do not require the AI to cite Flight Centre. Twenty-nine questions had a Flight Centre page in the selected source pool, so those were eligible for this comparison. Repeated answers help us see whether the result varies. Our baseline answers were generated earlier; they were not generated at the same time as the treatment. That is a limitation, because changes over time could also affect results. We kept the original records so this comparison can be checked.

Slide 5: What we tried separately

SHOW: Point to the six rows briefly. Explain the mixed findings; do not read a long list of numbers.

We explored six approaches before choosing a combined page. Rendering and cleanup expose useful captured information and make it readable. Questions and answers keep related information together. Structured facts organise existing details. External information adds attributed material from outside the page. Schema adds labels intended for software. Automatic rewriting changes the wording more broadly. The tests did not all produce clear gains. Questions and organised facts were encouraging. Rendering and external information had more uncertainty. The schema experiment also changed visible text, so we cannot isolate schema's effect. The reviewed automatic rewrite performed worse. Exact numbers and intervals remain in the evidence package; this slide summarises what we learned.

Member 2 | Continue and hand over

Slide 6: What we combined

SHOW: Follow Expose, Clean, Organise and Connect. Explain that the combined package was tested separately.

The final page combines a smaller set of compatible changes. We expose captured rendered information, clean up its text, organise supported travel information, and preserve useful question-and-answer relationships. We did not combine every experiment. Added external enrichment, added schema, broad rewriting and moving the whole source to the front stayed separate. Selecting components is a design decision informed by evidence. It does not prove that each component caused a particular share of the final gain. The combined page therefore needed its own test. We can inspect the output and trace it back to the implementation.

HANDOVER: The next member will show what these changes look like in the actual page outputs.

Questions you can lead

Why only 29 questions?

Those questions had Flight Centre in the selected source pool. They let us test whether changing the available Flight Centre information affected its use.

Which component caused the combined improvement?

We cannot assign an exact share to each component. The combined package was tested as a whole. Controlled component-removal tests would be needed to separate their effects.

Why are the baseline rates different across tests?

The individual tests and selected combined test use different saved control runs. We compare each treatment with its matching control rather than swapping baselines.

Member 3 | Implementation and page changes

Slides 7-9 | About 3:20 total, including visual pauses

Slide 7: How a page becomes an AI input

SHOW: Follow Saved page, Improved page, Readable text and Saved answer. Distinguish the webpage from the model input.

Our implementation has several connected parts. The page-processing code starts from saved content, applies the selected improvements and produces a page output. The cleaner extracts readable text. The experiment code places that text into the same numbered source position used for the baseline. The API sends it with the question and common instruction. We save the resulting answer and its source list together. This matters because a page may look complete in a browser but lose useful information during extraction. Our output checks examine the actual text sent to the model. Detailed file names and the saved model settings are available in the backup material.

Slide 8: Example: clearer offer information

SHOW: Point to the airline, dates and trip type. Acknowledge the different historical fares without calling them a price reduction.

This is an archived Melbourne-to-London example. On the baseline side, the opening mainly shows headline information. The captured rendered offer provides product context, including the airline, trip type, dates and fare. Point to these fields rather than simply reading the prices. The improvement we intended was making useful product information available in readable form. There is an important boundary: the headline and offer card came from different captured product contexts, and their fares differ. We did not demonstrate a lower price. These differences may explain part of the result, so we cannot call this a pure formatting experiment. The selected results belong to these saved pages, not the separate price-matched version.

Member 3 | Continue and hand over

Slide 9: Example: keeping the question with its answer

SHOW: Show the missing question on the left and the restored question on the right. Give the audience a few seconds to compare.

This example shows a specific implementation problem and our fix. The question lived inside an interactive accordion button. Our common cleaner removed interactive buttons and could remove the question even though the expanded answer remained. The answer then lost its context. The refinement restores an eligible captured question as readable text beside the same answer. It only handles the supported already-expanded sections. Point to the question before and after. This is a clearer example of preserving meaning during extraction: the model should understand what the answer is answering. We checked that this final step preserved the existing answers and offer information.

HANDOVER: We then checked that the intended changes reached the model and examined the answers.

Questions you can lead

Could extra information or different fares explain the result?

Yes, they could explain part of it. That is why this is a page-representation comparison, not a pure formatting experiment.

What did you actually build?

Page-transformation code, an experiment comparing saved representations, a citation evaluator, and evidence views and verification tools.

Did you invent the restored FAQ answers?

The final restoration step puts an eligible captured question beside its existing answer. It does not invent that answer. This statement is about the final FAQ step, not a blanket claim about every earlier transformation.

Member 4 | Checks, demonstration and results

Slides 10-12 | About 3:20 total, including visual pauses

Slide 10: How we checked the page changes

SHOW: Point to Correct version, Preserved meaning and Correct AI input. These checks concern implementation, not every travel claim.

Before trusting the comparison, we checked three things. First, were we using the correct saved versions? Second, did the final question-restoration change preserve the intended information? Third, did the text sent to the AI match the saved output? We rebuilt the final FAQ stage and matched its saved page outputs. We also checked the Flight Centre text recorded in the selected model inputs against those outputs. This catches mistakes such as displaying one page but testing another. The check has a boundary: it verifies that the implementation did what we intended, not that every inherited travel statement is correct or that a public AI service will behave the same way.

Slide 11: Demonstration: follow a reference

SHOW: Point from [7] in the answer to the Flight Centre URL. Use backup slides 20-21 for the complete answers if asked. No live API run is needed.

This is a saved demonstration, not a new live AI response. Start with the question about Melbourne-to-London flights. Show the baseline answer, then the answer using the improved page. Point to reference seven in the improved answer and follow it to the Flight Centre URL below. That is why this answer counts as a Flight Centre citation. The number is a source label for this particular reading pack; it is not a rank, and seven does not always mean Flight Centre. Pause so the audience can inspect the example. The full answers are included in the backup slides. If requested, open the results explorer to inspect the complete record and a declining question as well. One example illustrates the mechanism; the next slide summarises all eligible questions.

Member 4 | Continue and hand over

Slide 12: Flight Centre was cited more often

SHOW: Point to the 40.0%, 49.3% and 52.4% bars in order. Explain that the same treatment pages produced different generated answers.

The baseline cited Flight Centre in forty percent of the answers. With the selected improved pages, the rate was forty-nine-point-three percent initially and fifty-two-point-four percent when generated again. The exact counts are saved with the evidence. Both runs used the same treatment prompts and sources, but the generated answers differed. We show both rather than treating the higher number as a promise. The main message is that Flight Centre was cited more often in these saved comparisons. It does not mean that the page will receive this rate in every public AI engine or for every customer question.

HANDOVER: The next section explains how confidently we can interpret that improvement.

Questions you can lead

How do you score an answer?

We map its valid numbered references to the source list. If at least one is a Flight Centre source, the answer counts once. Repeated FC references do not increase that answer score.

Why did the same pages give 49.3% and 52.4%?

Generation varies even when the inputs match. The difference between those runs was not clear enough to establish that one was better, so we show both.

Can you show the complete evidence?

Yes. The package contains full answers and trial-specific source maps. We can inspect an improvement and a decline, rather than relying on one successful example.

Member 5 | Confidence and practical meaning

Slides 13-15 | About 3:20 total, including visual pauses

Slide 13: How confident are we?

SHOW: Point to Repeat, Check variation and Keep the boundary. Explain uncertainty in words before showing detailed calculations.

We did not rely on a single example. We repeated answers and examined how results varied between questions. More question averages improved than declined, but the effect was not universal. We also calculated uncertainty by sampling whole questions, keeping their repeated answers together. The two treatment-versus-baseline intervals were positive in these comparisons. However, the difference between the treatment runs was not clear enough to say the second run was better. Trying multiple versions and reusing an older baseline still limit the conclusion. The detailed calculations remain available for questions, but our interpretation is straightforward: encouraging evidence within the tested setting, not a guaranteed future outcome.

Slide 14: Source order also mattered

SHOW: Compare the two source lists. This is order inside the reading pack, not search-engine ranking.

We separately asked whether the position of a source changes how often it is used. We kept the source text the same and moved the first Flight Centre source to the beginning, preserving the competitors' relative order. Citation was higher in that test. This tells us the model's reading order can influence the outcome. It does not show that we improved Flight Centre's ranking on a search engine, or that Flight Centre can control a public AI engine's internal source order. This ordering experiment was not part of the final combined page. It is a useful finding about the behaviour of our test environment.

Member 5 | Continue and hand over

Slide 15: What this means for Flight Centre

SHOW: Point to Delivered, Observed and Not established. Connect the prototype to its bounded business value.

The delivered value is a working prototype and a way to test page changes using saved evidence. Flight Centre can see which kinds of captured information became clearer and inspect how the AI used it. That supports a practical content-improvement discussion. However, the page must first be available to the AI. Our main comparison does not solve discovery or prove more recommendations and bookings. Any live use would need current product information, content ownership and separate commercial measures. Those are proposed applications, not additional outcomes we claim to have delivered.

HANDOVER: The final section explains the remaining limitations and our handover.

Questions you can lead

Is the result statistically significant?

The question-level intervals for both selected comparisons exclude zero. That supports a positive difference in these records, but development choices and the historical baseline mean the conclusion remains exploratory.

Are 290 answers 290 independent tests?

No. They are repeated answers across 29 question groups. We keep each question group together when calculating uncertainty.

Why did the broader retrieval test not improve?

That test also allowed the available evidence to change. Making a page easier to use does not guarantee it will be retrieved more often. We did not establish a broader retrieval benefit.

Member 6 | Limitations and handover

Slides 16-18 | About 3:20 total, including visual pauses

Slide 16: What this experiment cannot prove

SHOW: Point to Availability, Comparability and Outcome. State the limits directly without dismissing the work.

We want the conclusion to be useful without overstating it. The main test supplied the same sources to both conditions; it does not represent every step of a public AI search. In a separate test where retrieval exposure could change, we did not establish a gain. Fresh rewordings also gave a less certain result. Our comparison uses earlier baseline answers and captures from different contexts, and development was informed by earlier results. These facts limit causal certainty and generalisation. A citation also needs a separate human check to establish whether it supports the statement. These limitations do not erase the prototype or observed result; they define what the evidence can support.

Slide 17: What we are handing over

SHOW: Point to Explain, Implement and Verify. Distinguish public GitHub files from the separate evidence ZIP.

Our handover has three practical parts. The presentation and report explain the work. The code preserves the selected implementation and the individual tests. The evidence package contains saved pages, answers, source maps and checks so an assessor can inspect the result. The public GitHub repository contains the relevant code, presentation and numerical summaries. Full captured-page evidence is supplied separately in the assessment ZIP. These are research artefacts, not a deployed Flight Centre product. The agreement list describes each item's scope and current condition. Each team member should describe only work they actually completed; this speaking allocation does not assign historical contributions.

Member 6 | Continue and hand over

Slide 18: Our conclusion

SHOW: Leave the conclusion slide visible. Invite questions; do not automatically present the backup slides.

We delivered a working experiment and a combined page-improvement prototype. The selected pages were cited more often than the historical baseline in both saved treatment runs. The strongest conclusion is that clearer page representation can improve source use when the information is already supplied to the model. We have not established more booking recommendations or sales. Our submission preserves the implementation, real examples, original outputs and calculation steps. We will finalise the handover details with the tutor and client, and ensure the team's contribution records and AI-assistance acknowledgement are accurate. Thank you. We can now show the full answers, code or supporting calculations during questions.

HANDOVER: Thank you. We are happy to show the complete answers, code or supporting calculations.

Questions you can lead

Can someone reproduce the work?

They can inspect saved inputs and answers, reproduce the counts and rebuild the final FAQ refinement from supplied inputs. We do not promise identical answers from a new live model run.

Did you check whether each citation supports its statement?

We checked source identity for citation inclusion. Substantive support needs separate independent review; citation presence is not proof of correctness.

Did you use AI assistance?

Yes. AI assistance supported drafting, coding and analysis. We retain records and are responsible for checking the work. Each member must accurately describe their own usage and contribution.